

Coding with Claude Code

The disciplined engineering loop for computational-materials research — seven phases, the constraint theory underneath them, and one units bug that only a human caught

Stephen Kerr · Computational Materials Science · Ontario Tech University

TITLE & HOOK — NOT VIBE-CODING · ~5 MIN

Coding with Claude Code, for people who validate the physics

- Not vibe-coding — the disciplined loop that produces code you can trust
- Most 'AI for coding' is written for web apps
- Workflow transfers cleanly; vocabulary doesn't
- You're fluent in the science — the new layer is the engineering workflow
- One running example, two threads, meeting at one number

Who this is for

- Comp-materials scientists & grad students who run atomistic/multiscale sims
- KMC, MD/statics, DFT — and write the code that does it
- Given: fluency in the science
- New: the AI-assisted engineering workflow on top
- Explicitly not for vibe coders

The two threads that converge

- Thread 1 — substrate: periodic_system_generator → Ni/Ni-Fe/Ni-Cr FCC(001) slabs, vacancies, step edges → LAMMPS+EAM minimize → pyKMC
- Thread 2 — correctness: ~23-commit harmonic-TST (Vineyard) prefactor campaign in pyKMC
- They meet at the Ni(100) surface-hop barrier dE
- Same dE the slabs study = same dE that pins the units regression test
- That convergence is why this is THE example

⊕ **anchor** — Thread 1 slab generation (Ni/Ni-Fe 95/5/Ni-Cr 95/5 FCC(001)) + Thread 2 HTST prefactor work, meeting at the Ni(100) surface-hop barrier dE.

THE REFRAME — DEV-SPEAK $\Leftrightarrow$ RESEARCH · ~7 MIN

The Rosetta stone: one substitution drives everything

- Deliverable: behaviour a user sees → a physical claim that survives scrutiny
- Every translation is a consequence of that one substitution
- PRD = study/method spec (a mini pre-registration)
- Bug = wrong PHYSICS — no crash, green suite
- QA = validating the physics, not click-testing a UI

The unit of work, translated

- Idea → a physics question, an artifact, or a code capability
- Feature → a new physical capability or measured quantity (the Vineyard backend)
- Bug → wrong physics, often green-and-silent (nu0 in Hz where the clock wants ps^{-1})
- Refactor → restructure the solver without changing the numbers
- Sprint → a simulation campaign (the ~23-commit HTST push)

⊕ **anchor** — Idea: 'give pyKMC physically-correct per-event rate prefactors instead of one fixed constant.' Feature: `pykmc/rate_constant/backends/htst.py`. Bug: nu0 Hz vs ps^{-1} , every rate 10^{12} too large, silently.

Artifacts, agents, and what NOT to carry over

- PRD: spec the END STATE, let implementation move underneath
- research.md asset: a sprint-lived cache allowed to rot
- Memory-less agent: why knowledge lives in contracts, not your head
- Trigger surface: the always-on cost of a skill — minimize it
- Don't import: 'ship fast, user won't notice' — our failure mode is silently-wrong

⊕ **anchor** — PRD example: 'Deliver D_vac(Ni,500K) with an Arrhenius barrier validated against an EAM-NEB value; archive trajkmc.xyz + reference_table.pickle.'
RateConstantConfig is the typed end-state contract.

REFERENCE · DESCEND OR SKIP

The dictionary — dev-speak $\rightleftharpoons$ research

software dev...	...materials research
Deliverable = behaviour a user sees	A physical claim, or a reusable capability, that...
PRD (product requirements document)	A study/method spec — a mini pre-registration...
Ticket / Jira card	One scoped sub-task of a study or code change — for...
Bug (wrong output / crash)	Wrong PHYSICS — often with no crash and a green...
Feature (new user-visible behaviour)	A new physical capability or a measured quantity...
Refactor	Restructure a solver/driver without changing the...
Sprint	A simulation campaign or a focused push to land one...
QA / click-testing a UI	Validating the physics: barrier vs NEB, detailed...

software dev...	...materials research
TDD / regression test	A test that pins a PHYSICAL contract and must be...
Vertical slice / tracer bullet	One end-to-end short run touching every pipeline...
software dev...	...materials research
Kanban board	The dependency graph of a study — stages in series...
Ralph loop / coding-agent execution...	An unattended execution loop over independent...
Token	A metered second compute budget; structurally, one...
Context budget / smart zone / dumb...	Finite (~100K usable), degrading attention; the...
Compaction vs handoff	Compaction buys continuity (lossy by omission you...
Skill / trigger surface	A reusable auto-triggering capability whose...
Memory-less agent	A fresh agent with no memory between tasks — which...
research.md asset	A throwaway cache of method details scoped to one...

software dev...	...materials research
Deep module / thin interface	The same design taste applied to solvers — the...
Green and wrong	A passing suite that has encoded the physics error...

FOUNDATIONS — WHY THE WORKFLOW HAS THIS SHAPE · ~9 MIN

The constraints that force the procedure

- Smart zone / dumb zone — usable attention ends ~100K tokens
- Bigger window doesn't move the line — just ships more dumb zone
- Feedback loops are the ceiling — not the model, not the prompt
- Tokens are a metered second compute budget — and mirror the KMC engine
- A procedure you don't understand is one you'll abandon when rushed

Smart zone / dumb zone — a correctness control

- Attention is quadratic: each token adds a row AND a column
- Past ~100K tokens the model makes measurably worse decisions
- For us the dumb zone is where silently-wrong physics is manufactured
- Sizing a task to fit the smart zone is a correctness control, not productivity advice
- Clear, don't compact — the Memento reset is reproducible; compaction carries sediment


⊕ **anchor** — pykmc/htst/constants.py: `omega_eV_to_hz()` and `hz_to_per_ps()` are two functions because they are two facts — the fine-grained correctness the agent holds in the smart zone and drops in the dumb zone.

THE PICTURE

Smart zone / dumb zone — a correctness control


Usable attention degrades before the advertised window

Decision quality vs. context size — the smart zone is a correctness control, not a capacity number



why it degrades

attention is quadratic



each token adds a row **and** a column — pairs grow as

$$\binom{n}{2} = n(n-1)/2$$

Sizing a task to fit the smart zone is a correctness control.

Clear, don't compact — start each unit of work from a fresh, small context.

Feedback loops are the ceiling — and ours can be green-and-wrong

- Web feedback: types, unit tests, lint, a screenshot
- Ours adds a second tier — the physics: detailed balance, NEB match, MSD linearity, one imaginary mode, ~THz prefactor
- 'Green and wrong': a passing suite that has ENCODED the error and defends it
- Deep modules → testable → strong feedback loops → AI does good work
- Bad codebase makes a bad agent — the 'in' of garbage-in is your architecture

⊕ **anchor** — HTST_Test_Report/run_htst_ab.py: paired constant-vs-htst runs on matched seeds, built to fail if HTST yields no finite nu0 — a physics-tier feedback loop. RateConstantConfig.require_one_negative_mode encodes physics in the type.

Tokens are the currency — and the KMC engine in miniature

- Tokens: a metered second budget, billed per call, input priced apart from output
- Feed a summary contract plus a file path, never the multi-MB dump
- `reference_table.pickle` IS a learned vocabulary; the KMC loop IS encode/compute/decode
- `n_current` = vocabulary hit (free); `n_new` = miss (`nsearch=20` ARTn search)
- Decode silently corrupts when the number carries a different unit than the consumer assumes


⊕ **anchor** — validation_basin: `trajkmc.xyz` >1MB, `reference_table.pickle` ~5.7MB, `lammps.log.0` ~11MB. The Ref Event is a token id; the `Hz→ps-1` bug is a decode failure.

THE PICTURE

Tokens are the currency — and the KMC engine in miniature

Same machine, twice over

An LLM token pipeline laid directly on the pyKMC main loop — the units bug lives in **DECODE**



THE SEVEN PHASES — THE SPINE OF THE METHOD · ~11 MIN

Idea → Research → Prototype → PRD → Kanban → Execution → QA


- Two conditional phases: Research (2) and Prototype (3) — skipped when small/familiar
- Last three (Kanban→Execution→QA) form a loop until the science converges
- QA failures become new tickets → back to Kanban
- A map, not a gate — tiny work skips 2, 3, most of 4
- Value: names for the steps, so nobody silently drops QA

THE PICTURE

The shape: a linear front half, a converging back-half loop

The seven-phase loop

Linear front half, two conditional phases, and a QA convergence loop



MATERIALS MEANING

1 Idea sharp dissatisfaction or a target number

2 Research cache method details; stamp build / date

3 Prototype sweep knobs till the catalog stabilizes

4 PRD RateConstantConfig + DO NOT MIX units header

5 Kanban dependency graph; cut vertically

6 Execution build launches real compute (mpirun / srun)

7 QA a human reads units & magnitudes

green tests are not validated physics

ONE LINE PER PHASE

The seven phases, in a simulation campaign

#	Phase	What it means for a study
1	Idea	A physics question, an artifact, or a code capability. State a sharp dissatisfaction or a...
2	Research (conditional)	Cache method details you'll forget — parameter semantics, the units convention, failure...
3	Prototype (conditional)	A throwaway convergence or recipe shakedown: sweep knobs on a tiny cell until the catalog...
4	PRD	A study/method spec — a mini pre-registration: system, conditions, observable + acceptance...
5	Kanban	The dependency graph of the study: stages in series (build→minimize→search→select→analyze)...
6	Execution	Same loop, but 'build' launches real compute (mpirun/srun) — genuinely asynchronous. One...
7	QA	Validating the physics. The agent-producible half is real (unit tests, type-check, lint...

Phase 1–3: Idea, Research, Prototype

- Idea: state a sharp dissatisfaction or a target number, not a solution
 - 'Rates use a constant prefactor and that's wrong' — not 'add a function vineyard()'
 - Research (conditional): cache method details into a sprint-lived research.md — stamp with conditions
 - Profiling harness, not a notebook: eskm $\sim 2.5\text{--}6\times$ faster than Python-FD, agreement 0.014%
 - Prototype (conditional): sweep knobs on a tiny cell, commit the winner, delete the losers
- ⊕ **anchor** — Phase 2: commit d52f541 profiling harness (eskm dynamical_matrix vs Python-FD) → benchmarks/htst_prefactor/RESULTS.md. Phase 3: in.minimize_slab_Ni_*.Imp variants; the NiAlH_jea default failed for Ni-Cr, FeNiCrCoAl-heaweight.setfl won.

Phase 4: PRD — spec the end state, grill the design

- PRD = study/method spec: system, conditions, observable + acceptance bands, output files
- NOT the implementation — over-pinned code rots against fast-moving scripts
- Grill the design: which potential? which seeds? what defines convergence? what's the null result?
- `require_one_negative_mode` encodes the physics in the type
- Record deferred scope: RpaBackend bare-Vineyard until the Sharia–Henkelman κ lands — deferred, not broken

⊕ **anchor** — `RateConstantConfig(style: constant|htst|rpa, free_radius, fd_step, nu0_min/max_THz, require_one_negative_mode)`. vineyard.py header: 'Units convention (DO NOT MIX) ... nu0 output: Hz (linear, NOT angular)'. Review commit a752002 (Hugo Moison).

THE CODE

Phase 4: PRD — spec the end state, grill the design

```
1 RateConstantConfig(  
2     style = "constant" | "htst" | "rpa", # thin interface literal  
3     free_radius = ...,  
4     fd_step = ...,  
5     nu0_min_THz = ..., nu0_max_THz = ...,  
6     require_one_negative_mode = True, # physics invariant AT the  
    interface:  
7 ) # a true first-order saddle has  
8     # exactly one imaginary mode.
```

Phase 5: Kanban — vertical slices, tickets are worktrees

- Board = dependency graph: stages in series, seeds/temps/compositions in parallel
- First ticket is a tracer bullet — one thin slice through every stage
- Cut vertically (build→minimize→search→select→analyze), never horizontally
- Tickets are literally worktree + branch names
- Independent fixes (freeze-core, units) ride parallel branches

⊕ **anchor** — HTST chain: 9327519 base → Vineyard+config+PHONON install → manager-ops → batched compute → reference backfill → active/refined backfill → recycling skip. Worktrees: pyKMC-htst-work, pyKMC-freezecore-work (fix/freeze-core-pbc), pyKMC-htst-units-work (fix/htst-rate-units).

PHASE 6 ZOOMED IN — THE EXECUTION LOOP · ~7 MIN

Ralph: a dumb bash loop that beats orchestration — and you already own one

- Backstop, one feature per iteration, append-only memory, per-iteration commit
- You already run it: `run_overnight.sh` under `nohup`
- 'done' must be runnable physics — not a green test suite
- Backstop is wall-clock + queue depth, not just iterations
- `set -e` vs `set -uo pipefail` IS the AFK-vs-attended decision

The 'done' predicate must be runnable physics

- Web: agent flips passes:true on green type-check + tests — defensible there
- Physics: a green suite tells you the code RAN, never that the physics is RIGHT
- passes = a machine-checkable acceptance band in eV/Å/ps, run by the real solver
- Predicate 1: require_one_negative_mode read off the actual dynamical matrix
- Predicate 2: finite-nu0 A/B gate; coverage<1.0 is a logged WARNING, not a pass

⊕ **anchor** — run_htst_ab.py: status='passed' only if the failures list (built from physics observables) is empty. A 511-atom Si smoke run hit 33.3% nu0 coverage — logged as partial, not a benchmark.

THE CODE

The 'done' predicate must be runnable physics

```
1 # run_htst_ab.py – status == "passed" only if failures is empty
2 if n_events <= 0:
3     failures.append("no reference events")
4 if "nu0" not in htst_table.columns:
5     failures.append("HTST table has no nu0 column")
6 if n_finite_nu0 <= 0:
7     failures.append("HTST produced no finite nu0 values")
8 if nu0_coverage < 1.0:
9     warnings.append(...) # WARNING, not a failure
```

Memory, sizing, commits, and the role shift

- Append, never update: a memory-less agent reconstructs state from the progress file
- Record the physics state: 'coverage 0.71; three events fell back to k0'
- One feature per iteration: compute_prefactors_batch shards the giant Hessian
- Per-iteration commit = bisectable; the surgical one-line units fix d705932 in its own worktree
- Ralph demotes you from anal-retentive planner to requirements-gatherer

⊕ **anchor** — Commit 82e7793 compute_prefactors_batch (batch size chosen by the d52f541 profiling harness); c9a74ad never-recompute-recycled = idempotency; chain ec6aa5f → ... → c9a74ad bisectable.

CONTEXT & HANDOFF — SURVIVING MANY FRESH AGENTS · ~4 MIN

Compaction buys continuity; handoff buys purity

- Grep the convergence line, not the log — 30 tokens vs 30,000
- Compaction is lossy by omission you don't control — drops the load-bearing caveat
- Handoff: a disposable temp-dir pointer doc carrying one task's slice
- Point, not copy — references stay correct as targets evolve
- Portable markdown → adversarial physics review by a different agent

The DIY sub-agent round trip

- Hand off a hard-to-judge bit to a prototype session — separate budget
 - It does the disposable work, returns a compressed verdict
 - RESULTS.md IS the handoff-back doc — stamped with its conditions
 - Same shape as a dispatched sub-agent, zero machinery
 - Promote durable facts to CONTEXT.md; the handoff is the courier, not the archive
- ⊕ **anchor** — Commit d52f541 round trip: the eskm-vs-Python-FD question needed a measurement; benchmarks/htst_prefactor/RESULTS.md returned the verdict (eskm faster, conditions attached) and the planner chose the eskm path.

QA, TAKEAWAYS & ATTRIBUTION · ~2 MIN

Phase 7: the units bug — and the one thing to remember

- ν_0 emitted in Hz (~1–10 THz), consumed as ps^{-1} → every rate 10^{12} too large
- Nothing crashed; the suite was green; the tests had pinned the error
- Fix d705932: one-line `hz_to_per_ps()` at the backend boundary
- Regression d6b921f: red by 10^{12} before, green after — at the Ni(100) dE where both threads meet
- QA = validating the physics. Green tests are not validated physics.

Attribution & the open questions

- Structure: five Matt Pocock talks; materials mapping is original to this collection
- Documents skills you may already run: grilling, to-prd, to-issues, tdd
- Open: merge Research + Prototype into one 'de-risk the method' phase?
- Open: keep the name 'Ralph' or 'coding-agent execution loop'?
- First cut — the seven-phase mapping is a proposal for discussion, not a settled standard

Speaker notes

